



Founder Advisory Board

Three advisors. One question. Deliberate disagreement.

THE PROBLEM

Founders make irreversible decisions with reversible-quality advice

- 01 Pricing, GTM, hiring — set wrong, they compound for years.
- 02 A real advisory board is slow, expensive, and unavailable pre-seed.
- 03 So founders ask an LLM instead.

THE INSIGHT

Ask an AI for advice and you get the average of the internet

Generic

Not conditioned on your stage or your numbers.

Unconditional

One answer, no "unless."

Confident

No visible assumptions to check.

And multi-agent "debate" doesn't fix it: **same model + same prompt + same evidence converges into expensive agreement.**

WHAT A REAL BOARD GIVES YOU

The disagreement is the product

It shows you **which assumption your decision rests on** — not the persona, not the vote count.

THE PRODUCT

A board that convenes on demand and argues over live evidence

"How should we price our AI project tracker?"

Three advisors take **independent positions**, then **cross-examine** each other.

The Chair synthesizes — **conditionally**, with **dissent preserved**.

You get a decision, its load-bearing assumptions, and a plan to test them.

THE BENCH

Three advisors who cannot agree by construction

VALUE PRICER

"You are almost certainly undercharging. Price the value, not the cost."

Enterprise, annual contracts, high ACV.
Skeptical of free tiers.

PLG ADVOCATE

"Friction is death. The product must sell itself before anyone talks to sales."

Free tier, self-serve, land-and-expand.
Skeptical of "contact sales" at seed.


UNIT-ECONOMICS HAWK

"A price is a claim about CAC payback and gross margin. Show me the arithmetic."

Takes no ideological side. Names the missing number and the threshold that flips it.

THE PIPELINE

Scrape → index → argue → synthesize



Independent-first is the load-bearing choice: advisors write in parallel and never see each other's drafts, so the disagreement survives to critique.

THE OUTPUT

Conditional advice, with the assumptions on the surface

Load-bearing assumptions

if it holds → do X · if it fails → invert

Position scorecard

every advisor, five named dimensions, no overall average

Dissent

the strongest surviving objection — a required field

Validation plan

the cheapest real-world tests that settle it

LATENCY HUD — MEASURED LIVE

moss retrieval <ms>

- think <s>
- voice <s>

Retrieval in milliseconds against model latency
in seconds — printed on every answer.



Most AI advice tells you what to do.
This tells you what you'd have to believe.

Structurally opposed advisors.

Grounded in live evidence.

Dissent preserved, tests attached.